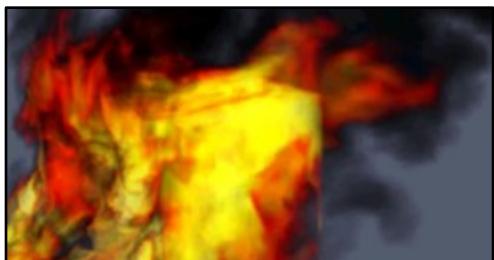




*Exceptional service in the national interest*



$$\partial a \cdots J_{a,\sigma^2}(\xi_1) = \frac{(\xi_1 - a)}{\sigma^2} f_{a,\sigma^2}(\xi_1);$$
$$\int_{\mathbb{R}_+} T(x) \cdot \frac{\partial}{\partial \theta} f(x, \theta) dx = M \left( T(\xi) \cdot \frac{\partial}{\partial \theta} \ln L \right)$$



# The Kokkos C++ Performance Portability EcoSystem

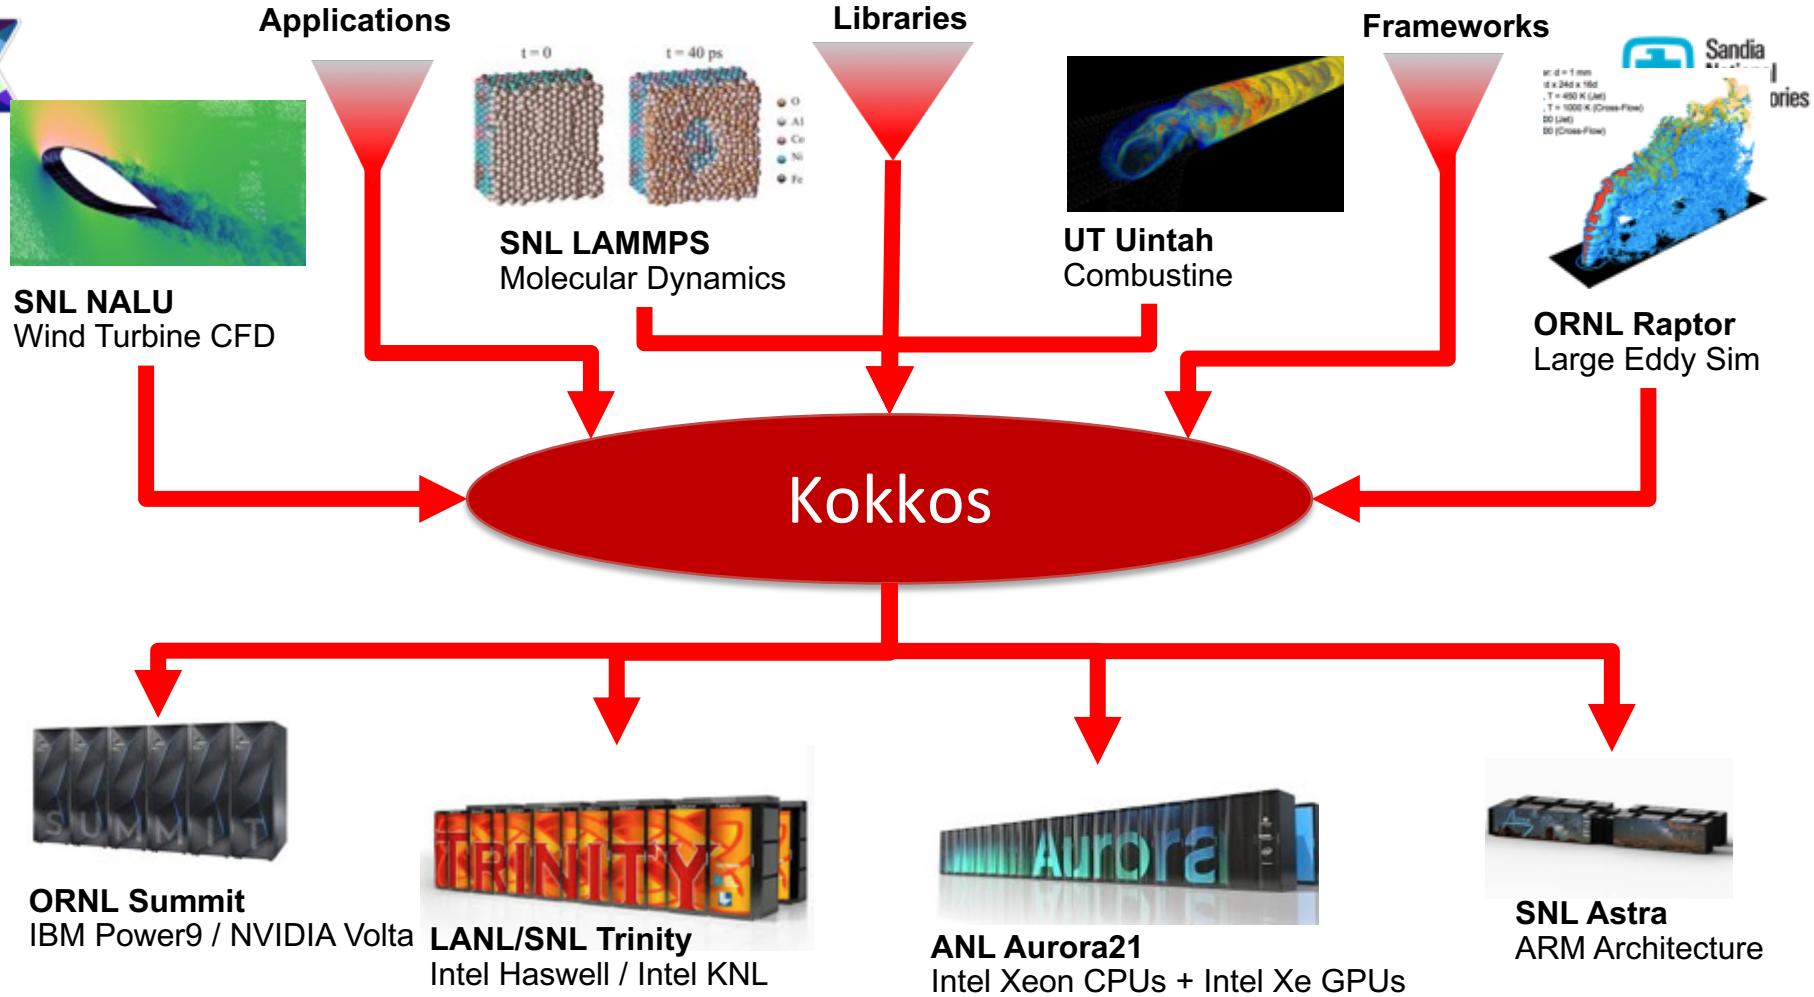
Unclassified Unlimited Release

*D. Sunderland, N. Ellingwood, D. Ibanez, S. Bova,  
J. Miles, D. Hollman, V. Dang*



***Christian R. Trott***, - Center for Computing Research  
Sandia National Laboratories/NM

Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA-0003525.



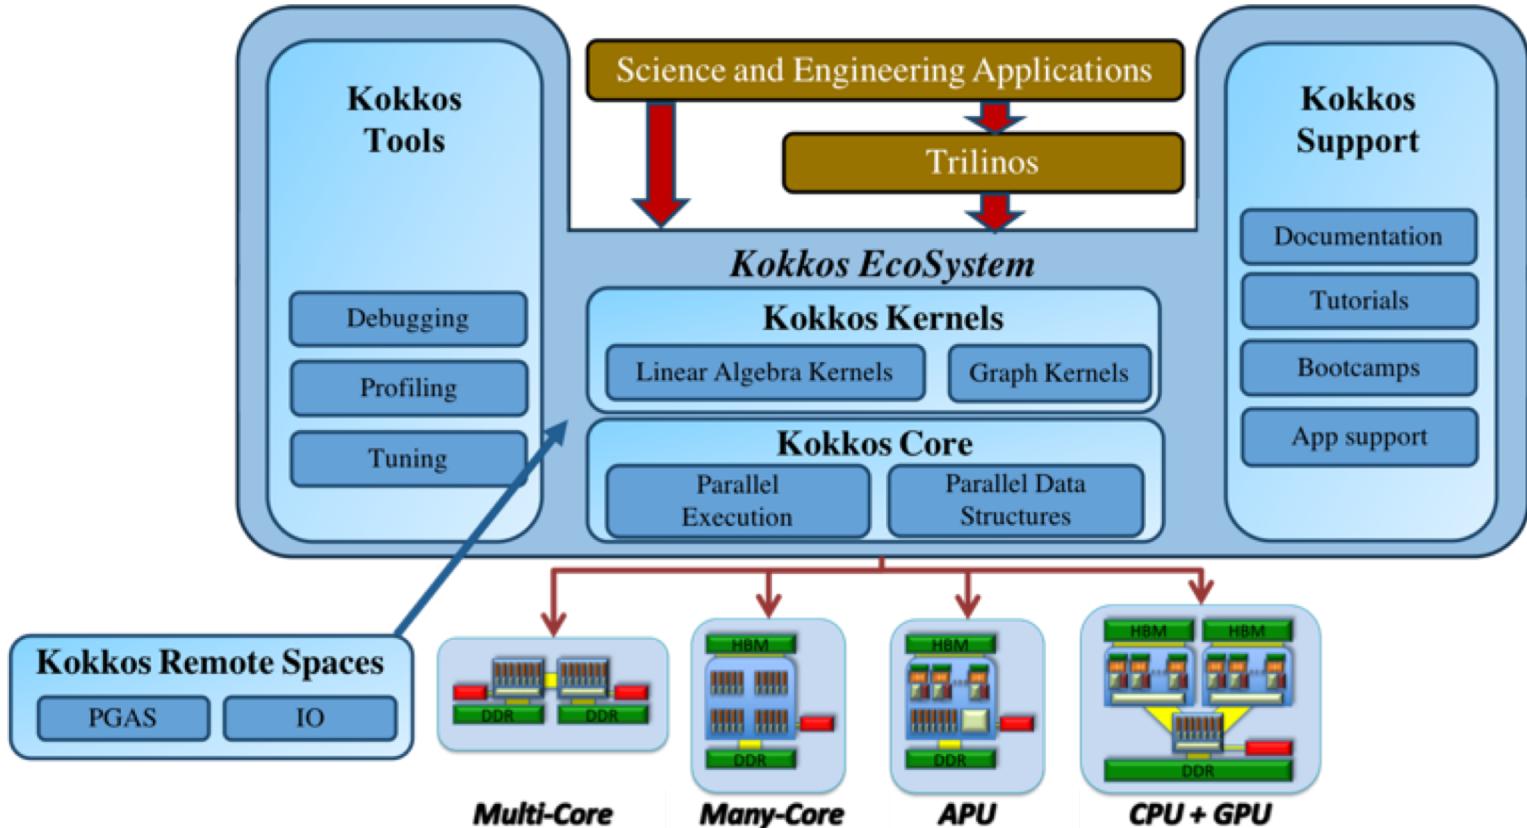


# Goals For Performance Portability



- One coherent approach to low level HPC performance portability needs
  - Parallel Execution
  - Data Structures and Management
  - Math Kernels
  - Tools
- Limit cognitive overload
  - Orthogonalization of concerns
  - Most of the time no explicit reference to backends (e.g. CUDA, or OpenMP)
- Off ramp via standards integration to limit scope
  - Invest into C++ standards work to make Kokkos a “sliding window” of advanced capabilities

# Kokkos EcoSystem





# Kokkos Development Team



- Dedicated team with a number of staff working most of their time on Kokkos
  - Main development team at Sandia in CCR – Sandia Apps are customers

**Kokkos Core:**

*C.R. Trott, D. Sunderland, N. Ellingwood, D. Ibanez, S. Bova, J. Miles, D. Hollman, V. Dang,  
soon: H. Finkel, N. Liber, D. Lebrun-Grandie, A. Prokopenko  
former: H.C. Edwards, D. Labreche, G. Mackey*

**Kokkos Kernels:**

*S. Rajamanickam, N. Ellingwood, K. Kim, C.R. Trott, V. Dang, L. Berger,*

**Kokkos Tools:**

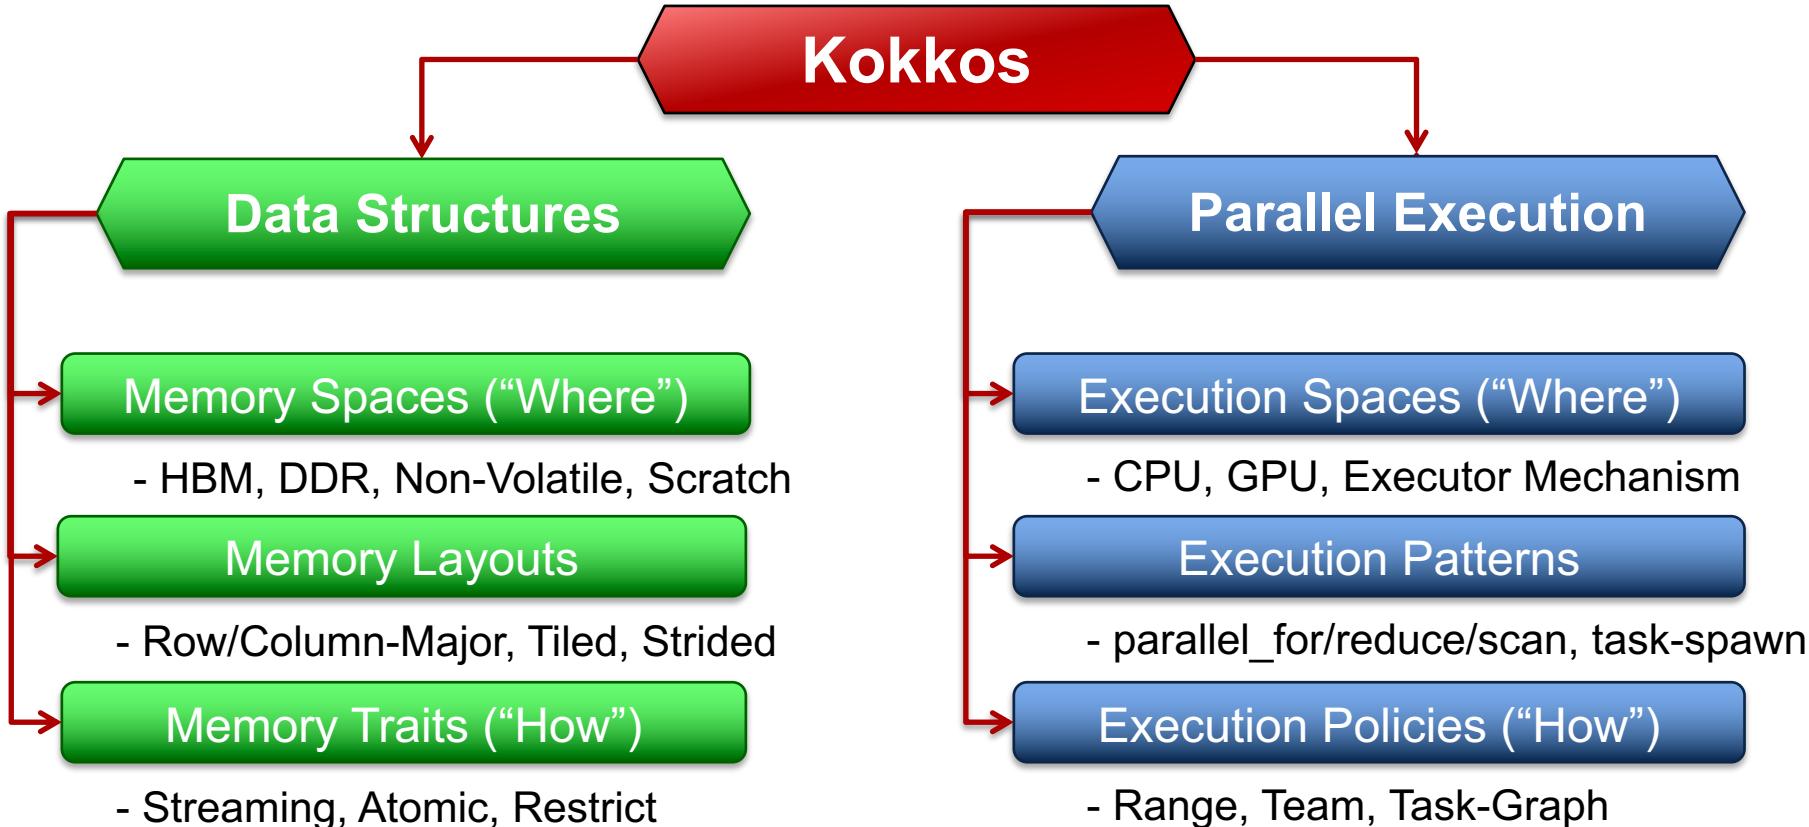
*S. Hammond, C.R. Trott, D. Ibanez, S. Moore*

**Kokkos Support:**

*C.R. Trott, G. Shipman, G. Lopez, G. Womeldorf,  
former: H.C. Edwards, D. Labreche, Fernanda Foertter*



# Kokkos Core Abstractions





# Patterns and Policy



- Reduce cognitive overload by reusing the same code structure

- **Parallel\_Pattern( ExecutionPolicy , FunctionObject [, ReductionArgs] )**

```
// Basic parallel for:  
parallel_for( N, Lambda );  
// Parallel for with dynamic scheduling:  
parallel_for( RangePolicy<Schedule<Dynamic>>(0,N), Lambda );  
// Parallel Reduce with teams:  
parallel_reduce( TeamPolicy<>(N,AUTO), Lambda, Reducer );  
// Parallel Scan with a nested policy  
parallel_scan( ThreadVectorRange(team_handle,N), Lambda );  
// Restriction pattern equivalent to #pragma omp single  
single( PerTeam(team_handle), Lambda );  
// Task Spawn  
task_spawn( TeamTask(scheduler, dependency), Task );
```

- Orthogonalize further via “require” mechanism to customize exec policy

```
auto exec_policy_low_latency = require(exec_policy,KernelProperty::HintLightWeight);
```



# Kokkos Core Capabilities



| Concept                  | Example   |
|--------------------------|---|
| Parallel Loops           | <code>parallel_for( N, KOKKOS_LAMBDA (int i) { ...BODY... });</code>  |
| Parallel Reduction       | <code>parallel_reduce( RangePolicy&lt;ExecSpace&gt;(0,N), KOKKOS_LAMBDA (int i, double&amp; upd) {<br/>    ...BODY...<br/>    upd += ...<br/>}, Sum&lt;&gt;(result));</code>  |
| Tightly Nested Loops     | <code>parallel_for(MDRangePolicy&lt;Rank&lt;3&gt; &gt; ({0,0,0},{N1,N2,N3},{T1,T2,T3},<br/>    KOKKOS_LAMBDA (int i, int j, int k) {...BODY...});</code>  |
| Non-Tightly Nested Loops | <code>parallel_for( TeamPolicy&lt;Schedule&lt;Dynamic&gt;&gt;( N, TS ), KOKKOS_LAMBDA (Team team) {<br/>    ... COMMON CODE 1 ...<br/>    parallel_for(TeamThreadRange( team, M(N)), [&amp;] (int j) { ... INNER BODY... });<br/>    ... COMMON CODE 2 ...<br/>});</code> |
| Task Dag                 | <code>task_spawn( TaskTeam( scheduler , priority), KOKKOS_LAMBDA (Team team) { ... BODY });</code>  |
| Data Allocation          | <code>View&lt;double**, Layout, MemSpace&gt; a("A",N,M);</code>   |
| Data Transfer            | <code>deep_copy(a,b);</code>  |
| Atomics                  | <code>atomic_add(&amp;a[i],5.0); View&lt;double*,MemoryTraits&lt;AtomicAccess&gt;&gt; a(); a(i)+=5.0;</code>  |
| Exec Spaces              | Serial, Threads, OpenMP, Cuda, HPX (experimental), ROCm (experimental)  |



# More Kokkos Capabilities

MemoryPool

DualView

Reducers

parallel\_scan

ScatterView

OffsetView

LayoutRight

StaticWorkGraph

RandomPool

sort

UnorderedMap

kokkos\_free

LayoutLeft

kokkos\_malloc

Vector

Bitset

LayoutStrided

ScratchSpace

ScratchSpace

ProfilingHooks



# Kokkos Kernels



- BLAS, Sparse and Graph Kernels on top of Kokkos and its View abstraction
  - Scalar type agnostic, e.g. works for any types with math operators
  - Layout and Memory Space aware
- Can call vendor libraries when available
- View have all their size and stride information => Interface is simpler

```
// BLAS                                // Kokkos Kernels
int M,N,K,LDA,LDB; double alpha, beta; double *A, *B, *C; double alpha, beta; View<double**> A,B,C;
dgemm('N','N',M,N,K,alpha,A,LDA,B,LDB,beta,C,LDC);           gemm('N','N',alpha,A,B,beta,C);
```

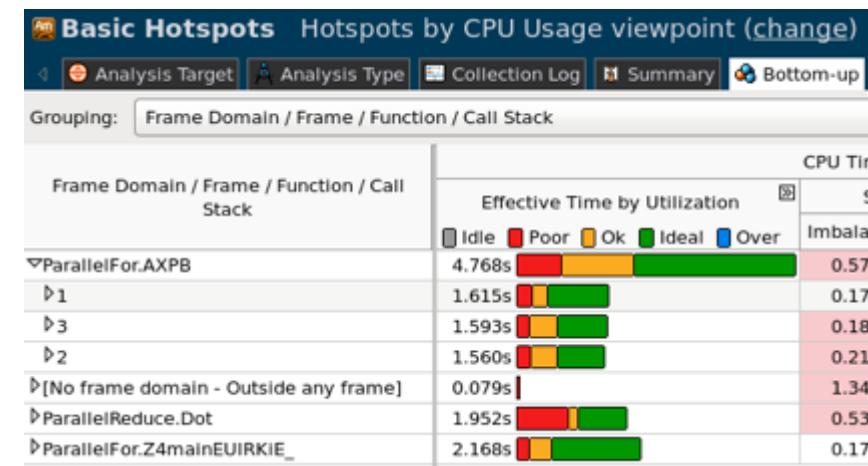
- Interface to call Kokkos Kernels at the teams level (e.g. in each CUDA-Block)

```
parallel_for("NestedBLAS", TeamPolicy<>(N,AUTO), KOKKOS_LAMBDA (const team_handle_t& team_handle) {
    // Allocate A, x and y in scratch memory (e.g. CUDA shared memory)
    // Call BLAS using parallelism in this team (e.g. CUDA block)
    gemv(team_handle,'N',alpha,A,x,beta,y)
});
```

# Kokkos-Tools Profiling & Debugging



- Performance tuning requires insight, but tools are different on each platform
- Insight into
- KokkosTools: Provide common set of basic tools + hooks for 3rd party tools
- One common issue abstraction layers obfuscate profiler output
  - Kokkos hooks for passing names on
  - Provide Kernel, Allocation and Region
- No need to recompile
  - Uses runtime hooks
  - Set via env variable





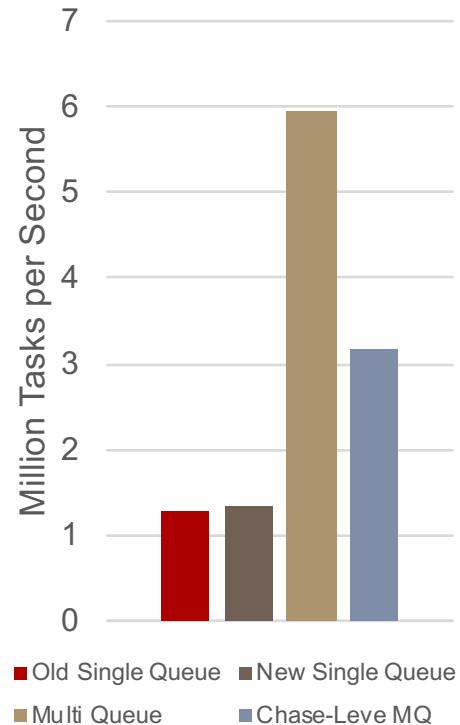
# Improved Fine Grained Tasking



- Generalization of TaskScheduler abstraction to allow user to be generic with respect to scheduling strategy and queue
- Implementation of new queues and scheduling strategies:
  - Single shared LIFO Queue (this was the old implementation)
  - Multiple shared LIFO Queues with LIFO work stealing
  - Chase-Lev minimal contention LIFO with tail (FIFO) stealing
  - Potentially more
- Reorganization of Task, Future, TaskQueue data structures to accommodate flexible requirements from the TaskScheduler
  - For instance, some scheduling strategies require additional storage in the Task

Questions: David Hollman

Fibonacci 30 (V100)



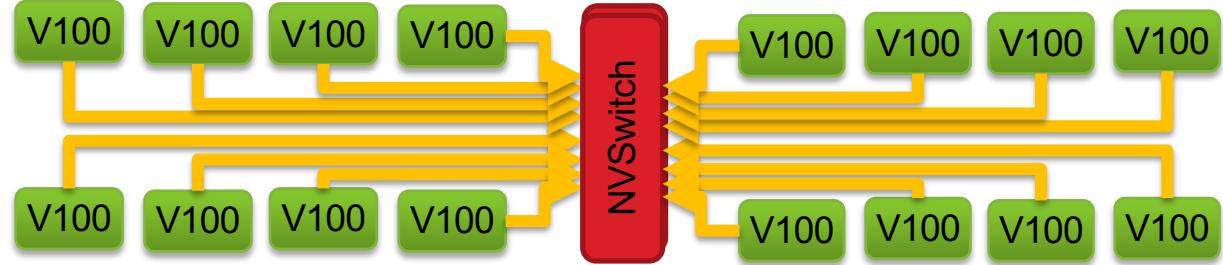


# Kokkos Remote Spaces: PGAS Support



- PGAS Models may become more viable for HPC with both changes in network architectures and the emergence of “super-node” architectures

- Example DGX2
- First “super-node”
- 300GB/s per GPU link



- Idea: Add new memory spaces which return data handles with shmem semantics to Kokkos View

- `View<double**[3], LayoutLeft, NVShmemSpace> a("A",N,M);`
- Operator `a(i,j,k)` returns:

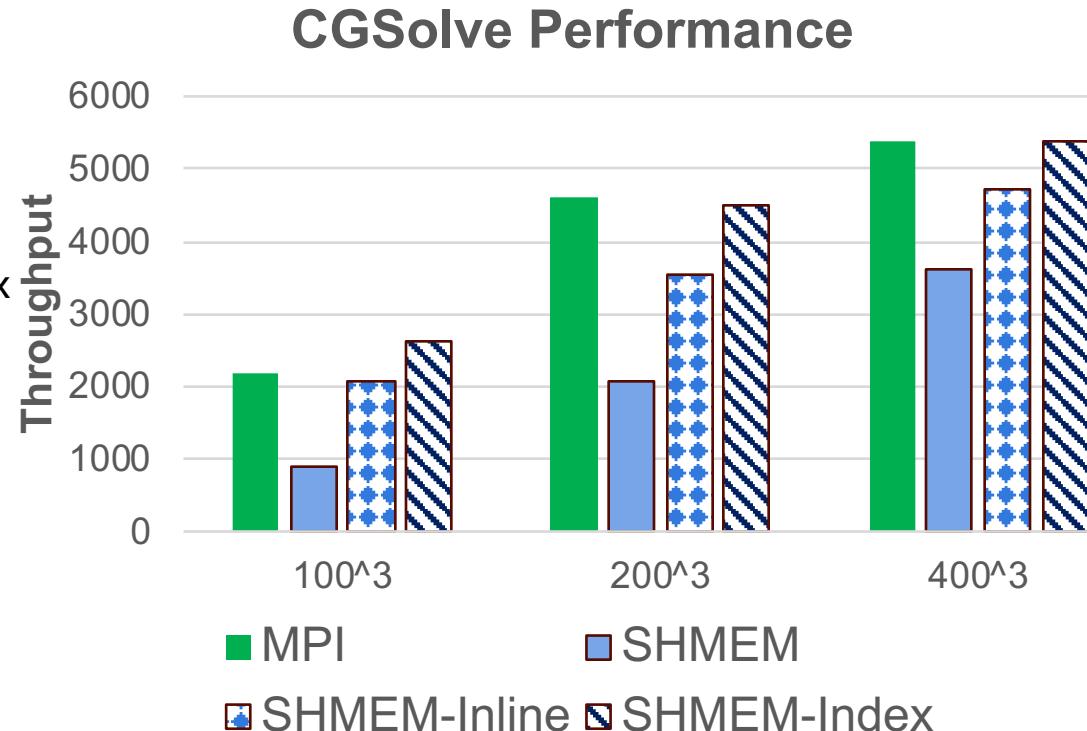
```
template<>
struct NVShmemElement<double> {
    NVShmemElement(int pe_, double* ptr_):pe(pe_),ptr(ptr_) {}
    int pe; double* ptr;
    void operator = (double val) { shmem_double_p(ptr,val,pe); }
};
```



# PGAS Performance Evaluation: miniFE



- Test Problem: CG-Solve
  - Using the miniFE problem N<sup>3</sup>
  - Compare to optimized CUDA
  - MPI version is using overlapping
  - DGX2 4 GPU workstation
  - Dominated by SpMV (Sparse Matrix Vector Multiply)
  - Make Vector distributed, and store global indices in Matrix
- 3 Variants
  - Full use of SHMEM
  - Inline functions by ptr mapping
    - Store 16 pointers in the View
  - Explicit by-rank indexing
    - Make vector 2D
    - Encode rank in column index



Warning: I don't think this is a viable thing in the next couple years for most of our apps!!



# Kokkos Based Projects



- Production Code Running Real Analysis Today
  - We got about **12** or so.
- Production Code or Library committed to using Kokkos and actively porting
  - Somewhere around **30**
- Packages In Large Collections (e.g. Tpetra, MueLu in Trilinos) committed to using Kokkos and actively porting
  - Somewhere around **50**
- Counting also proxy-apps and projects which are evaluating Kokkos (e.g. projects who attended boot camps and trainings).
  - Estimate **80-120** packages.

# Kokkos Users



Pacific Northwest  
NATIONAL LABORATORY



Rensselaer



TECHNISCHE  
UNIVERSITÄT  
MÜNCHEN

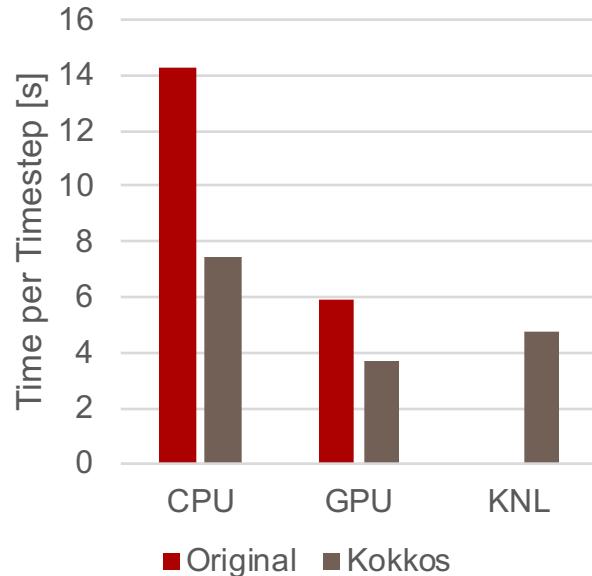


# Uintah

- System wide many task framework from University of Utah led by Martin Berzins
- Multiple applications for combustion/radiation simulation
- Structured AMR Mesh calculations
- Prior code existed for CPUs and GPUs
- Kokkos unifies implementation
- Improved performance due to constraints in Kokkos which encourage better coding practices

Questions: Dan Sunderlan

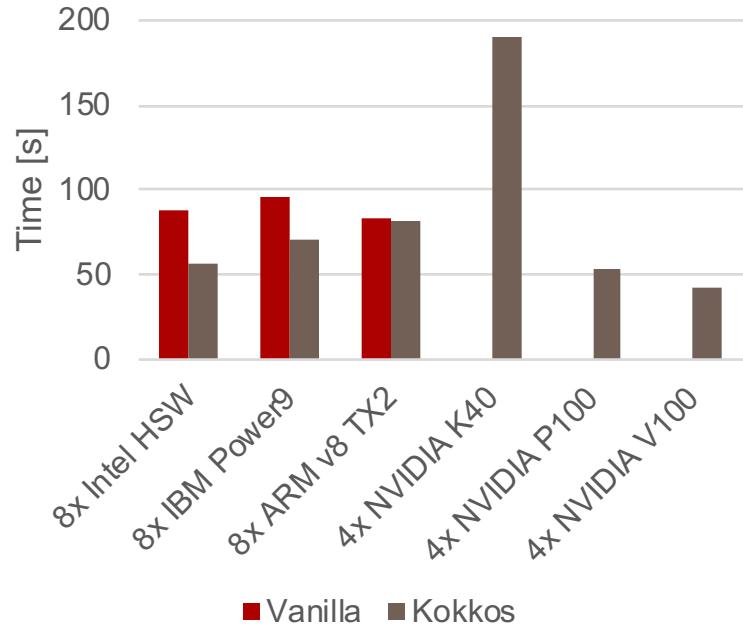
Reverse Monte Carlo  
Ray Tracing  $64^3$  cells





- Widely used Molecular Dynamics Simulations package
- Focused on Material Physics
- Over 500 physics modules
- Kokkos covers growing subset of those
- REAX is an important but very complex potential
  - USER-REAXC (Vanilla) more than 10,000 LOC
  - Kokkos version ~6,000 LOC
  - LJ in comparison: 200LOC
  - Used for shock simulations

Architecture Comparison  
Example in.reaxc.tatb /  
196k atoms / 100 steps

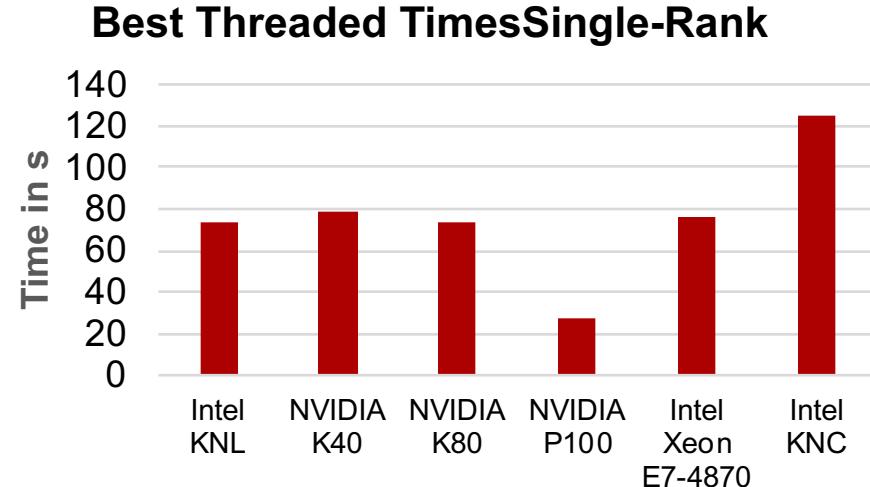
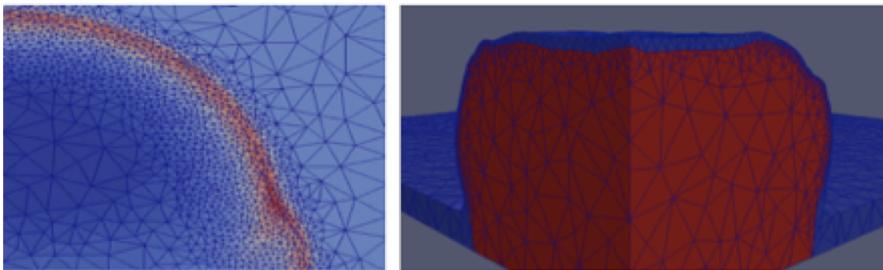




# Alexa

- Portably performant shock hydrodynamics application
- Solving multi-material problems for internal Sandia users
- Uses tetrahedral mesh adaptation

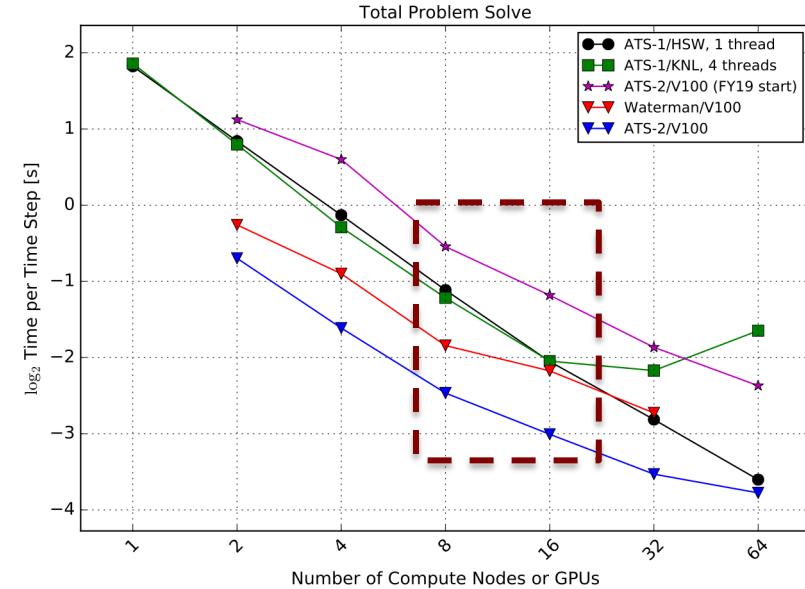
Questions: Dan Ibanez



- All operations are Kokkos-parallel
- Test case: metal foil expanding due to resistive heating from electrical current.



- Goal: solve aerodynamics problems for Sandia (transonic and hypersonic) on ‘leadership’ class supercomputers
- Solves compressible Navier-Stokes equations
- Perfect and reacting gas models
- Laminar and RANS turbulence models -> hybrid RANS-LES
- Primary discretization is cell-centered finite volume
- Research on high-order finite difference and discontinuous Galerkin discretizations
- Structured and unstructured grids



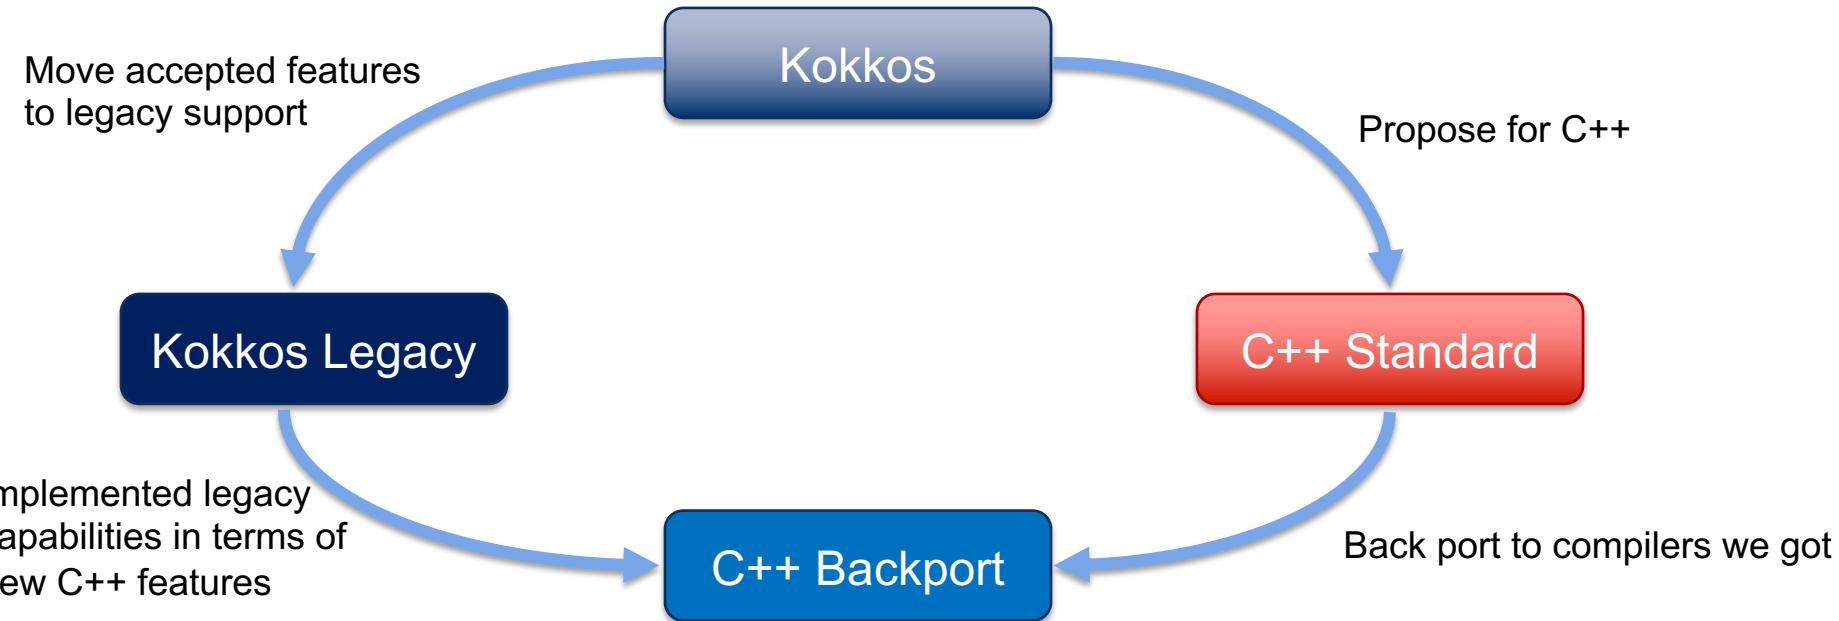
4 Sierra nodes (16x V100)  
equivalent to ~40 Trinity nodes  
(80x Haswell 16c CPU)



# Aligning Kokkos with the C++ Standard



- Long term goal: move capabilities from Kokkos into the ISO standard
  - Concentrate on facilities we really need to optimize with compiler





# C++ Features in the Works



- First success: **atomic\_ref<T>** in C++20
  - Provides atomics with all capabilities of atomics in Kokkos
  - **atomic\_ref(a[i])+=5.0;** instead of **atomic\_add(&a[i],5.0);**
- Next thing: **Kokkos::View => std::mdspan**
  - Provides customization points which allow all things we can do with **Kokkos::View**
  - Better design of internals though! => Easier to write custom layouts.
  - Also: arbitrary rank (until compiler crashes) and mixed compile/runtime ranks
  - We hope will land early in the cycle for C++23 (i.e. early in 2020)
- Also C++23: Executors and **Basic Linear Algebra** (just began design work)



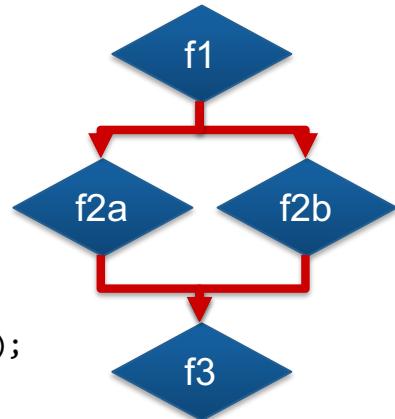
# Towards C++23 Executors

- C++ standard is moving towards more asynchronicity with Executors
  - Dispatch of parallel work consumes and returns new kind of future
- Aligning Kokkos with this development means:
  - Introduction of Execution space instances (CUDA streams work already)

```
DefaultExecutionSpace spaces[2];
partition( DefaultExecutionSpace(), 2, spaces);
// f1 and f2 are executed simultaneously
parallel_for( RangePolicy<>(spaces[0], 0, N), f1);
parallel_for( RangePolicy<>(spaces[1], 0, N), f2);
// wait for all work to finish
fence();
```

- Patterns return futures and Execution Policies consume them

```
auto fut_1 = parallel_for( RangePolicy<>("Funct1", 0, N), f1 );
auto fut_2a = parallel_for( RangePolicy<>("Funct2a", fut_1, 0, N), f2a);
auto fut_2b = parallel_for( RangePolicy<>("Funct2b", fut_1, 0, N), f2b);
auto fut_3 = parallel_for( RangePolicy<>("Funct3", all(fut_2a,fut2_b),0, N), f3);
fence(fut_3);
```





Sandia  
National  
Laboratories